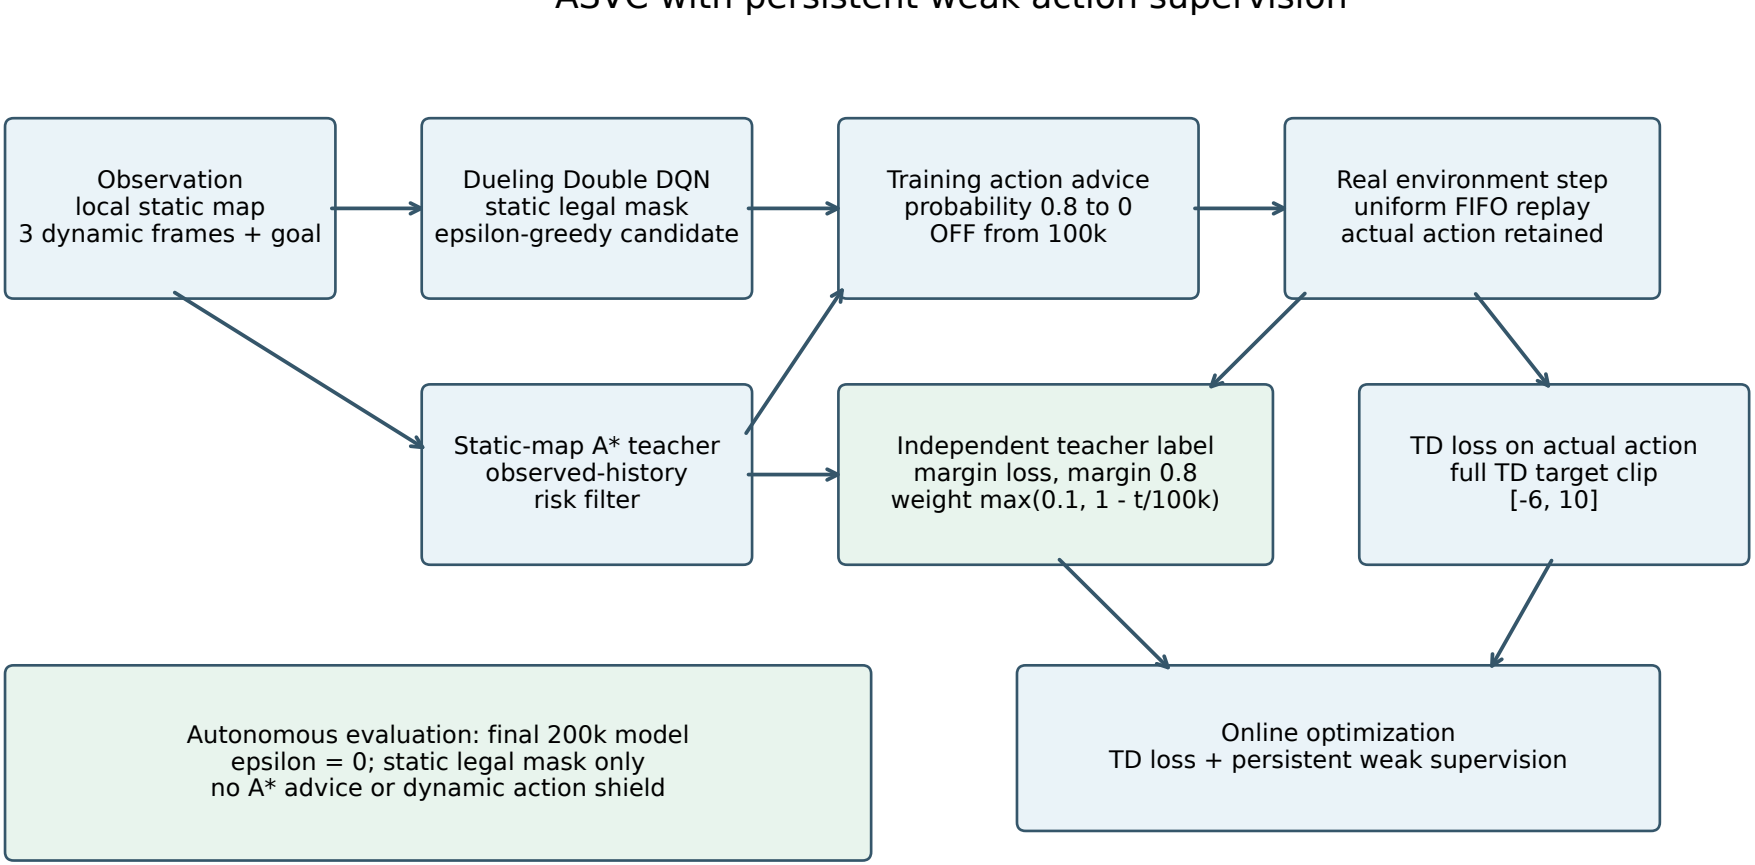
ASVC with persistent weak action supervision
Observation
local static map
3 dynamic frames + goal
Dueling Double DQN
static legal mask
epsilon-greedy candidate
Training action advice
probability 0.8 to 0
OFF from 100k
Real environment step
uniform FIFO replay
actual action retained
Static-map A* teacher
observed-history
risk filter
Independent teacher label
margin loss, margin 0.8
weight max(0.1, 1 - t/100k)
TD loss on actual action
full TD target clip
[-6, 10]
Autonomous evaluation: final 200k model
epsilon = 0; static legal mask only
no A* advice or dynamic action shield
Online optimization
TD loss + persistent weak supervision